

# END-TO-END SPEECH SUMMARIZATION USING RESTRICTED SELF-ATTENTION

*Roshan Sharma, Shruti Palaskar, Alan W Black and Florian Metze*

Carnegie Mellon University  
Pittsburgh PA, USA

roshansh@andrew.cmu.edu, spalaska@andrew.cmu.edu

## ABSTRACT

Speech summarization is typically performed by using a cascade of speech recognition and text summarization models. End-to-end modeling of speech summarization models is challenging due to memory and compute constraints arising from long input audio sequences. Recent work in document summarization has inspired methods to reduce the complexity of self-attentions, which enables transformer models to handle long sequences. In this work, we introduce a single model optimized end-to-end for speech summarization. We apply the restricted self-attention technique from text-based models to speech models to address the memory and compute constraints. We demonstrate that the proposed model learns to directly summarize speech for the How-2 corpus of instructional videos. The proposed end-to-end model outperforms the previously proposed cascaded model by 3 points absolute on ROUGE. Further, we consider the spoken language understanding task of predicting concepts from speech inputs and show that the proposed end-to-end model outperforms the cascade model by 4 points absolute F-1.

**Index Terms**— speech summarization, end-to-end, long sequence modeling, concept learning

## 1. INTRODUCTION

Summarization extracts and condenses desired information from the inputs, often text. Text can be summarized using abstraction or extraction [1]. Abstractive Text Summarization (ATS), generates a novel and concise summary of the input text. Abstractive summarization can be performed on multiple modalities [2, 3].

Speech Summarization is performed using a cascade of Automatic Speech Recognition (ASR) followed by Abstractive Text Summarization (ATS) [4, 5, 6]. [7] proposed an alternative cascade formulation- ASR followed by Concept Extraction and Summarization. They showed that specific and abstract concepts, extracted as nouns and noun-phrases, are useful intermediate representations for multimodal summarization. However, cascade architectures result in complex model structures with different modules optimized for different tasks, and errors in the ASR module degrade summariza-

tion performance. Therefore, we propose a single sequence model optimized end-to-end (E2E) for speech summarization.

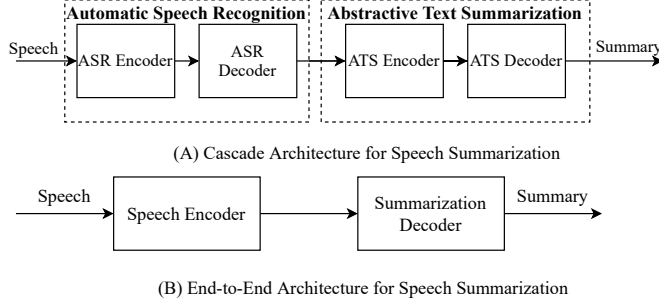
Speech summarization involves very long input sequences. The prohibitive quadratic computational cost of self-attention makes standard transformer models unsuitable for longer sequences. To address this, [8] uses segment-wise recurrence within transformer self-attention to provide longer context, and [9] compresses the segment level contexts and provides them as additional input to enable a longer context. Reformer [10] uses Locality Sensitive Hashing to compute localized self-attention in  $O(n \log n)$  and ETC [11] uses efficient global-local attention to scale to longer sequences. To reduce the complexity of self-attention to  $O(n)$ , Linformer [12] uses a low-rank factorization of the self-attention matrix, and Big Bird [13] uses a combination of sliding window, global and random attention. Longformer [14] uses different attention patterns for each layer and restricted dilated self-attention with task-specific global attention. These long sequence techniques have been evaluated on text inputs, where the input sequence lengths are often several hundred times smaller than sequence lengths of video-level speech (see Table 1).

Abstractive speech summarization uses the complete long sequence context to generate a summary. Long context ASR has been explored by training on longer sequences [15, 16, 17] or by passing context across utterances [18]. In this work, we (1) introduce a way to directly model speech summarization as an end-to-end task, (2) demonstrate the effectiveness of restricted self-attention for speech inputs, critical for the success of end-to-end speech summarization, and (3) show that such an end-to-end model can also be applied to learn concepts directly from speech inputs, a potential spoken language understanding task.

## 2. BACKGROUND

### 2.1. Cascade and E2E Modeling

Traditionally, speech summarization is modeled as a cascade of speech recognition and text summarization [4, 5]. Figure 1 shows the cascade and end-to-end approaches to speech summarization. The cascade approach benefits from strong ASR models pre-trained on large amounts of speech and summa-



**Fig. 1.** Speech Summarization: Cascade and End-to-End Model Architectures

ization models like BART [19] trained on large amounts of text data. However, errors in ASR are compounded due to the cascade architecture, which serves as motivation for direct end-to-end modeling.

## 2.2. Concept Learning

In [7], authors propose cascade multimodal speech summarization via semantic concept learning, where speech is transcribed, and then represented as a sequence of semantic concepts. These concepts are then input to a summarization model that generates abstractive summaries. Concepts are domain-specific noun phrases extracted automatically from the manually annotated summaries. Their cascade approach is able to generate summaries given good concept predictions (which are reliant on ASR predictions), and in this work, we evaluate the benefits of learning such concepts directly from long speech inputs as a language understanding task. Further, we model speech summarization as an end-to-end task optimized for summarization.

**Table 1.** Statistics of the How-2 2000h Dataset used for model training and evaluation. The mean and maximum statistics of N- the input length in frames, and L- the output length (in tokens) is shown.

Set	Max N	Mean N	Mean L	Max L
Train	145,082	9,806.58	60.54	173
Test	39,537	9,866.55	60.29	152

## 3. PROPOSED APPROACH

### 3.1. Restricted Self-Attentions

Different from other speech tasks like ASR, the speech inputs for summarization are much longer (5s for ASR versus 100s for summarization). Table 1 shows the average and maximum frame lengths of input speech, and output token lengths for summarization.

The high computational complexity makes it intractable to train video-level speech models on a GPU. Consider  $N$  is the length of the input speech sequence, and  $L$  is the length of the output token sequence ( $N \gg L$  from Table 1). It is known that encoder self-attention has a computational complexity of  $O(N^2)$ , decoder self-attention has a complexity of  $O(L^2)$ , and encoder-decoder source-target attention has a complexity of  $O(NL)$ . In order to make end-to-end training possible for summarization, the computational complexity of the encoder self-attention needs to be reduced. Inspired by [14, 20], we break down the self-attention computation into fixed sized context windows of size  $W$ . For each sequence element, a surrounding context of width  $W/2$  on each side is considered while computing the self-attention result. The number of such windows required will be  $P = N/W$ , and the cost of the encoder-self attention is now reduced to  $O(PW^2)$ , which is smaller than  $O(N^2)$ . To further reduce the computational complexity, we can drop one element for every  $D$  elements, i.e., dilation. Dilation further reduces the complexity to  $O(P(W/D)^2)$ .

### 3.2. End-to-End Speech Summarization

Given input speech frames for an *entire* video, we propose to directly summarize it into short, abstractive, textual summaries. The objective of mapping long speech frames (details in Table 1) onto significantly shorter textual tokens makes this an End-to-End Speech Summarization task. As training summarization models from scratch is challenging, we pretrain the sequence model using ASR. Then, the encoder-decoder model is fine-tuned for speech summarization.

### 3.3. End-to-End Concept Learning

Semantic concepts were shown to be a strong grounding aspect across modalities, especially to bridge the gaps in cascaded speech summarization [7]. Intermediate concept learning can be useful for controllability of generated summaries. Abstract concepts were extracted in [2] by transcribing the videos into text format, and then training a concept extractor. We contend that it would be useful to train a concept extractor from speech end-to-end. As we propose end-to-end speech summarization, we also evaluate the utility of our model to generate semantic concepts directly from speech. Given input speech at the video-level, we train our language understanding model to output a sequence of abstract semantic concepts.

## 4. EXPERIMENTAL SETUP

### 4.1. Dataset and Evaluation

The How-2 Dataset [21] contains 2000h of instructional videos with corresponding text transcripts, video, speech, translations, and summaries. Two tasks are evaluated: (a)

**Table 2.** Word Error Rate (WER) (%) for Test and Held Test sets of the 2000h How-to Corpus. Window Size of 20 is used for Restricted Self-Attention

Encoder	Decoder	Test WER (%)
Transformer	Transformer	10.2
Conformer	Transformer	<b>9.1</b>
+ Restricted Self-Attention	Transformer	9.3

Abstract Concept Generation from Speech, and (b) Abstractive Speech Summarization. Concept Generation is evaluated using Precision, Recall, and F-1 score. Summarization is evaluated using standard metrics ROUGE [22], METEOR [23], and BERTScore [24].

## 4.2. Model Details

ESPNet [25] is used for speech model training. Our conformer encoder uses 2-fold convolutional subsampling followed by 12 encoder layers with feed-forward dimension 2048, and 8 attention heads. The transformer decoder has 6 layers with feed-forward dimension 512 and 4 attention heads. ASR models are trained with joint Connectionist Temporal Classification (CTC)-Attention [26] with the weight for CTC training set to 0.3. The videos are trimmed to 100s for the video-level speech tasks owing to compute constraints. SpecAugment [27] is used during model training and fine-tuning. We use 40-dimension filterbank and 3-dimensional pitch features for training all models. Huggingface transformers [28] is used to fine-tune text-only cascade models. BART-large and BART-base [19] are fine-tuned on How2 transcript and summaries. Our code<sup>1</sup> and pre-trained models<sup>2</sup> have been released.

**Table 3.** Effect of Window Size and Dilation in Self-Attention of the Speech Encoder on E2E Summarization Model Training. W is the Window Size, and D is the dilation factor (Section 3 for details).

W	D	ROUGE-L	METEOR	BERTScore
20	<b>X</b>	52.0	26.5	90.5
40	<b>X</b>	<b>53.1</b>	<b>27.3</b>	<b>90.6</b>
60	<b>X</b>	52.5	27.1	90.5
100	5	51.9	26.3	90.5

## 5. RESULTS AND DISCUSSION

### 5.1. Speech Recognition

As described in Section 3, our speech summarization model is pre-trained for ASR. Table 2 shows the impact of encoder

type and attention type on Word Error Rate (WER). The use of a conformer [30] encoder improves ASR results by over 1 % absolute compared to the transformer whereas restricted self-attention results in a slight decrease in performance.

### 5.2. Speech Summarization

Table 4 highlights summarization results on three types of models: ground-truth text-based models (considered the topline scores), ASR-output based Cascade models, and direct E2E models. BART-large and BART-base [19] are fine-tuned on ASR predicted text( generated using the best ASR from 2) to establish the cascade baselines. BART-large outperforms BART-base in ROUGE, METEOR, and BERT Scores among the cascade models. Conformer ASR coupled with BART leads to strong cascade models that outperform previous works. The restricted self-attention based ASR model is then fine-tuned on the summarization data which results in the E2E summarization model. The E2E model outperforms the best cascade model on all metrics with 4x fewer parameters, indicating that the end-to-end model is able to produce more fluent, semantically relevant summaries.

Difference in METEOR between our models is correlated with difference in ROUGE-L scores. METEOR scores are content based, and missing out key noun phrases lowers the METEOR scores. From Table 4, it is clear that the cascade model and E2E models have lower METEOR Scores than the Cascade Concept Model and the Ground-truth models as the latter are better at retaining these noun phrases.

**Window Size and Dilation :** To understand the impact of context window size on summarization performance, we train models with different window sizes using a subset of the training data. From Table 3, a window size of  $W = 40$  seems to yield the best ROUGE-L scores, while a smaller window of  $W = 20$  yields a lower ROUGE-L score. An optimal window size is neither too short nor too long. Short windows lose important context, while longer windows incorporate less relevant context. From the first and last row, dilation reduces the computational complexity significantly while retaining comparable performance.

**Qualitative Examples :** Table 5 demonstrates two kinds of errors that we attribute to the cascade effect - missing content words(in blue), and mistranscribed words(in red). The proposed E2E approach mitigates the impact of these two types of errors, improving ROUGE and METEOR scores.

### 5.3. Concept Learning

Table 6 evaluates the end-to-end concept learning model. Concepts being non-sequential text, we evaluate on Precision, Recall, and F1. The baseline is a cascade of two modules- ASR and *predicted* Text2Concept model, and the proposed end-to-end Speech2Concept model outperforms the baseline by 4 points on F-1 and 10 points on Precision.

<sup>1</sup><https://github.com/espnet/espnet>

<sup>2</sup>[https://huggingface.co/espnet/roshanish.how2\\_asr\\_raw\\_ft\\_sum\\_valid.acc](https://huggingface.co/espnet/roshanish.how2_asr_raw_ft_sum_valid.acc)

**Table 4.** Summarization Performance of Topline, Cascade and E2E Models using automatic (ROUGE and METEOR) and semantic evaluation metrics (BERTScore).

	Model	Parameters	ROUGE-1	ROUGE-2	ROUGE-L	METEOR	BERTScore
Topline	Groundtruth Text						
	+ BART-base Summarization [3]	140M	<b>64.0</b>	46.4	58.9	31.7	-
Cascade	Conformer ASR	107M					
	+ BART-large Summarization	400M	59.2	38.8	52.3	27.8	90.6
	+ BART-base Summarization	140M	57.6	36.3	50.3	25.6	90.3
	S2S- PredText2Summary[2]	-	-	-	46.1	22.9	-
	ASR + BERTSum [29]	-	49.3	28.8	48.2	-	-
	Kaldi ASR + Concept2Summary [7]	-	-	-	51.4	<b>30.4</b>	-
E2E	Conformer Encoder						
	+ Transformer Decoder	<b>104M</b>	60.73	<b>44.9</b>	<b>56.10</b>	29.3	<b>91.53</b>

**Table 5.** Errors in the Cascade and E2E Approaches

E2E	DEFENDING AGAINST A SELF-DEFENSE TECHNIQUE IS THE PRINCIPLE OF THE ATTACKER'S ARM. <b>LEARN HOW TO STRIKE</b> AGAINST A SELF-DEFENSE IN THIS FREE VIDEO FROM AN <b>INDUCTEE IN THE US MARTIAL ARTS HALL OF FAME.</b>
	<b>DEF OR DEFANGING THE SNAKE</b> IS A SELF-DEFENSE TECHNIQUE THAT TAKES THE ATTACKER'S STRIKE OUT OF PLAY. <b>DEFANG THE SNAKE</b> WITH TIPS FROM A <b>MARTIAL ARTS INSTRUCTOR</b> IN THIS FREE VIDEO ON SELF DEFENSE.
Cascade	SELF DEFENSE TECHNIQUES MADE EASY ! <b>LEARN HOW TO STRIKE</b> AGAINST A PUNCH IN THIS FREE VIDEO FROM AN <b>INDUCTEE IN THE US MARTIAL ARTS HALL OF FAME .</b>
Ground Truth	

**Table 6.** Evaluation of Baseline and Proposed Concept Learning Models using Recall, Precision and F-1 Score

Model	Precision	Recall	F-1
Predicted Text2Concept [7]	52.5	<b>57.3</b>	54.8
Speech2Concept	<b>62.3</b>	55.8	<b>58.8</b>

## 6. CONCLUSION

In this paper, we model speech summarization as an end-to-end sequence task starting from video-level input speech to generate abstractive textual summaries as the output. We address the long speech input frames problem by applying restricted self-attention to help us achieve this task without running into severe memory and compute bottlenecks. Our approach at least outperforms a strong text-based summarization model, and at best, demonstrates strong performance compared to previous approaches to speech summarization (cascaded pipeline models). We also demonstrate the effects of various window sizes and dilations on summarization, concluding that optimal window sizes are neither too long nor too short. Using restricted self-attention and a conformer based speech recognizer, we achieve a competitive result on speech recognition on the commonly used How2 dataset. Finally, we demonstrate the potential of such end-to-end modeling on a Speech2Concept task that could be useful for downstream summarization as well as other speech-based tasks that ear-

lier represented speech by predicted text from an automatic speech recognizer.

## 7. REFERENCES

- [1] U. Hahn and I. Mani, "The challenges of automatic summarization," *Computer*, vol. 33, no. 11, pp. 29–36, 2000.
- [2] Shruti Palaskar, Jindřich Libovický, Spandana Gella, and Florian Metze, "Multimodal abstractive summarization for how2 videos," in *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, Florence, Italy, 2019, pp. 6587–6596, Association for Computational Linguistics.
- [3] Tiezheng Yu, Wenliang Dai, Zihan Liu, and Pascale Fung, "Vision guided generative pre-trained language models for multimodal abstractive summarization," 2021.
- [4] Chenguang Zhu, Ruochen Xu, Michael Zeng, and Xuedong Huang, "A hierarchical network for abstractive meeting summarization with cross-domain pretraining," in *Findings of the Association for Computational Linguistics: EMNLP 2020*, Online, 2020, pp. 194–203, Association for Computational Linguistics.
- [5] Dana Rezaadegan, Shlomo Berkovsky, Juan C Quiroz, A Baki Kocaballi, Ying Wang, Liliana Laranjo, and Enrico Coiera, "Automatic speech summarisation: A scoping review," *arXiv preprint arXiv:2008.11897*, 2020.
- [6] Potsawee Manakul, Mark Gales, and Linlin Wang, "Abstractive spoken document summarization using hierarchical model with multi-stage attention diversity optimization," 2020.
- [7] Shruti Palaskar, Ruslan Salakhutdinov, Alan W. Black, and Florian Metze, "Multimodal Speech Summarization Through Semantic Concept Learning," in *Proc. Interspeech 2021*, 2021, pp. 791–795.
- [8] Zihang Dai, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc Le, and Ruslan Salakhutdinov, "Transformer-XL: Attentive language models beyond a fixed-length context," in *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, Florence, Italy, 2019, pp. 2978–2988, Association for Computational Linguistics.

- [9] Jack W. Rae, Anna Potapenko, Siddhant M. Jayakumar, Chloe Hillier, and Timothy P. Lillicrap, “Compressive transformers for long-range sequence modelling,” in *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. 2020, OpenReview.net.
- [10] Nikita Kitaev, Lukasz Kaiser, and Anselm Levskaya, “Reformer: The efficient transformer,” in *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. 2020, OpenReview.net.
- [11] Joshua Ainslie, Santiago Ontanon, Chris Alberti, Vaclav Cvicek, Zachary Fisher, Philip Pham, Anirudh Ravula, Sumit Sanghai, Qifan Wang, and Li Yang, “ETC: Encoding long and structured inputs in transformers,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Online, 2020, pp. 268–284, Association for Computational Linguistics.
- [12] Sinong Wang, Belinda Li, Madian Khabsa, Han Fang, and Hao Ma, “Linformer: Self-attention with linear complexity,” *arXiv preprint arXiv:2006.04768*, 2020.
- [13] Manzil Zaheer, Guru Guruganesh, Kumar Avinava Dubey, Joshua Ainslie, Chris Alberti, Santiago Ontanon, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, and Amr Ahmed, “Big bird: Transformers for longer sequences,” in *Advances in Neural Information Processing Systems*, H. Larochelle, M. Ranzato, R. Hadsell, M. F. Balcan, and H. Lin, Eds. 2020, vol. 33, pp. 17283–17297, Curran Associates, Inc.
- [14] Iz Beltagy, Matthew E. Peters, and Arman Cohan, “Longformer: The long-document transformer,” *CoRR*, vol. abs/2004.05150, 2020.
- [15] Laurent El Shafey, Hagen Soltau, and Izhak Shafran, “Joint Speech Recognition and Speaker Diarization via Sequence Transduction,” in *Proc. Interspeech 2019*, 2019, pp. 396–400.
- [16] Takaaki Hori, Niko Moritz, Chiori Hori, and Jonathan Le Roux, “Transformer-based long-context end-to-end speech recognition,” in *Interspeech 2020, 21st Annual Conference of the International Speech Communication Association, Virtual Event, Shanghai, China, 25-29 October 2020*, Helen Meng, Bo Xu, and Thomas Fang Zheng, Eds. 2020, pp. 5011–5015, ISCA.
- [17] Takaaki Hori, Niko Moritz, Chiori Hori, and Jonathan Le Roux, “Advanced long-context end-to-end speech recognition using context-expanded transformers,” 2021.
- [18] Ryo Masumura, Naoki Makishima, Mana Ihori, Akihiko Takashima, Tomohiro Tanaka, and Shota Orihashi, “Hierarchical transformer-based large-context end-to-end asr with large-context knowledge distillation,” in *ICASSP 2021 - 2021 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2021, pp. 5879–5883.
- [19] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, and Luke Zettlemoyer, “BART: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension,” in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, Online, 2020, pp. 7871–7880, Association for Computational Linguistics.
- [20] Niko Moritz, Takaaki Hori, and Jonathan Le Roux, “Capturing multi-resolution context by dilated self-attention,” in *ICASSP 2021 - 2021 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2021, pp. 5869–5873.
- [21] Ramon Sanabria, Ozan Caglayan, Shruti Palaskar, Desmond Elliott, Loïc Barrault, Lucia Specia, and Florian Metze, “How2: a large-scale dataset for multimodal language understanding,” *arXiv preprint arXiv:1811.00347*, 2018.
- [22] Chin-Yew Lin, “ROUGE: A package for automatic evaluation of summaries,” in *Text Summarization Branches Out*, Barcelona, Spain, 2004, pp. 74–81, Association for Computational Linguistics.
- [23] Michael Denkowski and Alon Lavie, “Meteor universal: Language specific translation evaluation for any target language,” in *Proceedings of the Ninth Workshop on Statistical Machine Translation*, Baltimore, Maryland, USA, 2014, pp. 376–380, Association for Computational Linguistics.
- [24] Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi, “Bertscore: Evaluating text generation with BERT,” in *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. 2020, OpenReview.net.
- [25] Shinji Watanabe, Takaaki Hori, Shigeki Karita, Tomoki Hayashi, Jiro Nishitoba, Yuya Unno, Nelson Enrique Yalta Soplín, Jahn Heymann, Matthew Wiesner, Nanxin Chen, Adithya Renduchintala, and Tsubasa Ochiai, “ESPnet: End-to-end speech processing toolkit,” in *Proceedings of Interspeech*, 2018, pp. 2207–2211.
- [26] Suyoun Kim, Takaaki Hori, and Shinji Watanabe, “Joint ctc-attention based end-to-end speech recognition using multi-task learning,” in *2017 IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP 2017, New Orleans, LA, USA, March 5-9, 2017*. 2017, pp. 4835–4839, IEEE.
- [27] Daniel S. Park, William Chan, Yu Zhang, Chung-Cheng Chiu, Barret Zoph, Ekin Dogus Cubuk, and Quoc V. Le, “SpecAugment: A simple augmentation method for automatic speech recognition,” in *INTERSPEECH*, 2019.
- [28] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander Rush, “Transformers: State-of-the-art natural language processing,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, Online, 2020, pp. 38–45, Association for Computational Linguistics.
- [29] Takatomo Kano, Atsunori Ogawa, Marc Delcroix, and Shinji Watanabe, “Attention-based multi-hypothesis fusion for speech summarization,” 2021.
- [30] Anmol Gulati, Chung-Cheng Chiu, James Qin, Jiahui Yu, Niki Parmar, Ruoming Pang, Shibo Wang, Wei Han, Yonghui Wu, Yu Zhang, and Zhengdong Zhang, Eds., *Conformer: Convolution-augmented Transformer for Speech Recognition*, 2020.